

```
In [ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [ ]: !pip install lightning  
!pip install pytorch_forecasting
```

Requirement already satisfied: lightning in /usr/local/lib/python3.10/dist-packages (2.1.2)
Requirement already satisfied: PyYAML<8.0,>=5.4 in /usr/local/lib/python3.10/dist-packages (from lightning) (6.0.1)
Requirement already satisfied: fsspec[http]<2025.0,>2021.06.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (2023.6.0)
Requirement already satisfied: lightning-utilities<2.0,>=0.8.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (0.10.0)
Requirement already satisfied: numpy<3.0,>=1.17.2 in /usr/local/lib/python3.10/dist-packages (from lightning) (1.23.5)
Requirement already satisfied: packaging<25.0,>=20.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (23.2)
Requirement already satisfied: torch<4.0,>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (2.1.0+cu121)
Requirement already satisfied: torchmetrics<3.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (1.2.1)
Requirement already satisfied: tqdm<6.0,>=4.57.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (4.66.1)
Requirement already satisfied: typing-extensions<6.0,>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from lightning) (4.9.0)
Requirement already satisfied: pytorch-lightning in /usr/local/lib/python3.10/dist-packages (from lightning) (2.1.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from fsspec[http]<2025.0,>2021.06.0->lightning) (2.31.0)
Requirement already satisfied: aiohttp!=4.0.0a0,!4.0.0a1 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]<2025.0,>2021.06.0->lightning) (3.9.1)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from lightning-utilities<2.0,>=0.8.0->lightning) (67.7.2)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch<4.0,>=1.12.0->lightning) (3.13.1)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch<4.0,>=1.12.0->lightning) (1.12)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch<4.0,>=1.12.0->lightning) (3.2.1)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch<4.0,>=1.12.0->lightning) (3.1.2)
Requirement already satisfied: triton==2.1.0 in /usr/local/lib/python3.10/dist-packages (from torch<4.0,>=1.12.0->lightning) (2.1.0)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning) (23.1.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning) (6.0.4)
Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning) (1.9.4)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning) (1.4.0)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning) (1.3.1)
Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning) (4.0.3)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2->torch<4.0,>=1.12.0->lightning) (2.1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning) (2023.11.17)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch<4.0,>=1.12.0->lightning) (1.3.0)
Requirement already satisfied: pytorch_forecasting in /usr/local/lib/python3.10/dist-packages (1.0.0)
Requirement already satisfied: fastapi>=0.80 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (0.105.0)
Requirement already satisfied: lightning<3.0.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (2.1.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (3.7.1)
Requirement already satisfied: optuna<4.0.0,>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (3.5.0)
Requirement already satisfied: pandas<=3.0.0,>=1.3.0 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (1.5.3)
Requirement already satisfied: pytorch-optimizer<3.0.0,>=2.5.1 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (2.12.0)
Requirement already satisfied: scikit-learn<2.0,>=1.2 in /usr/local/lib/python3.10/dist-packages (from pyto

rch_forecasting) (1.2.2)
Requirement already satisfied: scipy<2.0,>=1.8 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (1.11.4)
Requirement already satisfied: statsmodels in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (0.14.0)
Requirement already satisfied: torch<3.0.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from pytorch_forecasting) (2.1.0+cu121)
Requirement already satisfied: anyio<4.0.0,>=3.7.1 in /usr/local/lib/python3.10/dist-packages (from fastapi>=0.80->pytorch_forecasting) (3.7.1)
Requirement already satisfied: pydantic!=1.8,!=1.8.1,!=2.0.0,!=2.0.1,!=2.1.0,<3.0.0,>=1.7.4 in /usr/local/lib/python3.10/dist-packages (from fastapi>=0.80->pytorch_forecasting) (1.10.13)
Requirement already satisfied: starlette<0.28.0,>=0.27.0 in /usr/local/lib/python3.10/dist-packages (from fastapi>=0.80->pytorch_forecasting) (0.27.0)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from fastapi>=0.80->pytorch_forecasting) (4.9.0)
Requirement already satisfied: PyYAML<8.0,>=5.4 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (6.0.1)
Requirement already satisfied: fsspec[http]<2025.0,>2021.06.0 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (2023.6.0)
Requirement already satisfied: lightning-utilities<2.0,>=0.8.0 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (0.10.0)
Requirement already satisfied: numpy<3.0,>=1.17.2 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (1.23.5)
Requirement already satisfied: packaging<25.0,>=20.0 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (23.2)
Requirement already satisfied: torchmetrics<3.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (1.2.1)
Requirement already satisfied: tqdm<6.0,>=4.57.0 in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (4.66.1)
Requirement already satisfied: pytorch-lightning in /usr/local/lib/python3.10/dist-packages (from lightning<3.0.0,>=2.0.0->pytorch_forecasting) (2.1.2)
Requirement already satisfied: alembic>=1.5.0 in /usr/local/lib/python3.10/dist-packages (from optuna<4.0.0,>=3.1.0->pytorch_forecasting) (1.13.0)
Requirement already satisfied: colorlog in /usr/local/lib/python3.10/dist-packages (from optuna<4.0.0,>=3.1.0->pytorch_forecasting) (6.8.0)
Requirement already satisfied: sqlalchemy>=1.3.0 in /usr/local/lib/python3.10/dist-packages (from optuna<4.0.0,>=3.1.0->pytorch_forecasting) (2.0.23)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas<=3.0.0,>=1.3.0->pytorch_forecasting) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas<=3.0.0,>=1.3.0->pytorch_forecasting) (2023.3.post1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn<2.0,>=1.2->pytorch_forecasting) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn<2.0,>=1.2->pytorch_forecasting) (3.2.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch<3.0.0,>=2.0.0->pytorch_forecasting) (3.13.1)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch<3.0.0,>=2.0.0->pytorch_forecasting) (1.12)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch<3.0.0,>=2.0.0->pytorch_forecasting) (3.2.1)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch<3.0.0,>=2.0.0->pytorch_forecasting) (3.1.2)
Requirement already satisfied: triton==2.1.0 in /usr/local/lib/python3.10/dist-packages (from torch<3.0.0,>=2.0.0->pytorch_forecasting) (2.1.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->pytorch_forecasting) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->pytorch_forecasting) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->pytorch_forecasting) (4.46.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->pytorch_forecasting) (1.4.5)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->pytorch_forecasting) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->pytorch_forecasting) (3.1.1)
Requirement already satisfied: patsy>=0.5.2 in /usr/local/lib/python3.10/dist-packages (from statsmodels->pytorch_forecasting) (0.5.4)
Requirement already satisfied: Mako in /usr/local/lib/python3.10/dist-packages (from alembic>=1.5.0->optuna<4.0.0,>=3.1.0->pytorch_forecasting) (1.3.0)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.10/dist-packages (from anyio<4.0.0,>=3.7.1->fastapi>=0.80->pytorch_forecasting) (3.6)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.10/dist-packages (from anyio<4.0.0,>=3.7.1->fastapi>=0.80->pytorch_forecasting) (1.3.0)

Requirement already satisfied: exceptiongroup in /usr/local/lib/python3.10/dist-packages (from anyio<4.0.0, >=3.7.1->fastapi==0.80->pytorch_forecasting) (1.2.0)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (2.31.0)

Requirement already satisfied: aiohttp!=4.0.0a0,!4.0.0a1 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (3.9.1)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from lightning-utilities<2.0,>=0.8.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (67.7.2)

Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy==0.5.2->statsmodels->pytorch_forecasting) (1.16.0)

Requirement already satisfied: greenlet!=0.4.17 in /usr/local/lib/python3.10/dist-packages (from sqlalchemy==1.3.0->optuna<4.0.0,>=3.1.0->pytorch_forecasting) (3.0.2)

Requirement already satisfied: MarkupSafe==2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch<3.0.0,>=2.0.0->pytorch_forecasting) (2.1.3)

Requirement already satisfied: mpmath==0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch<3.0.0,>=2.0.0->pytorch_forecasting) (1.3.0)

Requirement already satisfied: attrs==17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (23.1.0)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (6.0.4)

Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (1.9.4)

Requirement already satisfied: frozenlist==1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (1.4.0)

Requirement already satisfied: aiosignal==1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (1.3.1)

Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (4.0.3)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (3.3.2)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (2.0.7)

Requirement already satisfied: certifi==2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->fsspec[http]<2025.0,>2021.06.0->lightning<3.0.0,>=2.0.0->pytorch_forecasting) (2023.11.17)

```
In [ ]: import lightning.pytorch as pl
        from lightning.pytorch.callbacks import EarlyStopping, LearningRateMonitor
        from lightning.pytorch.loggers import TensorBoardLogger
        import numpy as np
        import pandas as pd
        import torch

        import optuna
        from optuna.integration import PyTorchLightningPruningCallback

        from pytorch_forecasting import Baseline, TemporalFusionTransformer, TimeSeriesDataSet
        from pytorch_forecasting.data import GroupNormalizer, TorchNormalizer, MultiNormalizer, EncoderNormalizer
        from pytorch_forecasting.metrics import MAE, SMAPE, PoissonLoss, QuantileLoss, MAPE, MultivariateNormalDist

        import pandas as pd
        import numpy as np

        #from pytorch_forecasting.models.temporal_fusion_transformer.tuning import optimize_hyperparameters
```

Creating Multi Step Dataset

```
In [ ]: df = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMH_Clean_UCI_Power_Consumpti
cpi = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YM_Clean_France_CPI_Electric
temp = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMH_Clean_TempISD.csv")
cci = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YM_Clean_CCI.csv")
gt1 = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMW_Clean_GoogleTrends1.csv")
gt2 = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMW_Clean_GoogleTrends2.csv")
gt3 = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMW_Clean_GoogleTrends3.csv")
```

```
<ipython-input-33-968a71a3a344>:1: DtypeWarning: Columns (2,3,4,5,6,7) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMH_Clean_UCI_Power_Consumption_Dataset.csv")
<ipython-input-33-968a71a3a344>:3: DtypeWarning: Columns (20) have mixed types. Specify dtype option on import or set low_memory=False.
temp = pd.read_csv("/content/drive/MyDrive/BDA696Project/Datasets/processsed/YMH_Clean_TempISD.csv")
```

```
In [ ]: ISD = pd.read_csv('/content/drive/MyDrive/BDA696Project/TFT_Model_Datasets/ISD_cleaned.csv')
ISD['Date'] = pd.to_datetime(ISD['New_Date'])
ISD=ISD[['Date', 'Wind_Speed_Norm', 'Prec_Depth_Norm', 'Air_Temp_Norm']]
# ISD = ISD.drop_duplicates(subset='New_Date', keep='first')
# ISD.sort_values(by='New_Date', ascending = False)
# Wind_Speed_Agg = ISD.resample('D', on='New_Date').Wind_Speed_Norm.mean()
# Prec_Depth_Agg = ISD.resample('D', on='New_Date').Prec_Depth_Norm.mean()
# Air_Temp_Agg = ISD.resample('D', on='New_Date').Air_Temp_Norm.mean()
```

```
<ipython-input-34-a0624b3daacb>:1: DtypeWarning: Columns (27) have mixed types. Specify dtype option on import or set low_memory=False.
ISD = pd.read_csv('/content/drive/MyDrive/BDA696Project/TFT_Model_Datasets/ISD_cleaned.csv')
```

```
In [ ]: # df = pd.read_csv("Downloads/processsed/YMH_Clean_UCI_Power_Consumption_Dataset.csv")
# cpi = pd.read_csv("Downloads/processsed/YM_Clean_France_CPI_Electricity.csv")
# temp = pd.read_csv("Downloads/processsed/YMH_Clean_TempISD.csv")
# cci = pd.read_csv("Downloads/processsed/YM_Clean_CCI.csv")
# gt1 = pd.read_csv("Downloads/processsed/YMW_Clean_GoogleTrends1.csv")
# gt2 = pd.read_csv("Downloads/processsed/YMW_Clean_GoogleTrends2.csv")
# gt3 = pd.read_csv("Downloads/processsed/YMW_Clean_GoogleTrends3.csv")
```

```
In [ ]: df['Date'] = pd.to_datetime(df['Date'])
df=df.query('(Date.dt.hour==0 or Date.dt.hour==6 or Date.dt.hour==12 or Date.dt.hour==18) and Date.dt.minute<30')
df=df.replace('?',0.0)
df.head()
```

```
Out [ ]:
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity	Sub_metering_1	Sub_
0	2007-01-01 00:00:00	00:00:00	2.580	0.136	241.970	10.600	0.000	
360	2007-01-01 06:00:00	06:00:00	2.460	0.064	241.130	10.200	0.000	
720	2007-01-01 12:00:00	12:00:00	2.478	0.000	235.300	10.400	0.000	
1080	2007-01-01 18:00:00	18:00:00	1.416	0.000	239.130	5.800	0.000	
1440	2007-01-02 00:00:00	00:00:00	0.442	0.122	241.060	1.800	0.000	

Creating and using function to foward fill weekly and monthly values

```
In [ ]: def fillblanks(df_temp):
temp = pd.DataFrame(pd.date_range(start=df_temp['Date'].min(), end= df_temp['Date'].max()),columns=['Date'])
temp['Date'] = pd.to_datetime(temp['Date'])
temp = temp.merge(df_temp, on='Date', how='left')
temp = temp.replace(0, np.nan).ffill()
return temp
```

```
In [ ]: #temp['Date'] = pd.to_datetime(cci['Date'])
#df=pd.merge_asof(df,temp, on='Date')

#df.dtypes

cpi['Date'] = pd.to_datetime(cpi['Date'])
df=df.merge(fillblanks(cpi), on='Date', how='left')
df['CPI'] = df['CPI'].replace(0, np.nan).ffill()
```

```

cci['Date'] = pd.to_datetime(cci['Date'])
df=df.merge(fillblanks(cci), on='Date', how='left')
df['CCI'] = df['CCI'].replace(0, np.nan).ffill()

ISD['Date'] = pd.to_datetime(ISD['Date'])
df=df.merge(fillblanks(ISD), on='Date', how='left')
df['Wind_Speed_Norm'] = df['Wind_Speed_Norm'].replace(0, np.nan).ffill()
df['Prec_Depth_Norm'] = df['Prec_Depth_Norm'].replace(0, np.nan).ffill()
df['Air_Temp_Norm'] = df['Air_Temp_Norm'].replace(0, np.nan).ffill()

gt1['Date'] = pd.to_datetime(gt1['Date'])
df=df.merge(fillblanks(gt1), on='Date', how='left')
df['GT1_Hits'] = df['GT1_Hits'].replace(0, np.nan).ffill()

gt2['Date'] = pd.to_datetime(gt2['Date'])
df=df.merge(fillblanks(gt2), on='Date', how='left')
df['GT2_Hits'] = df['GT2_Hits'].replace(0, np.nan).ffill()

gt3['Date'] = pd.to_datetime(gt3['Date'])
df=df.merge(fillblanks(gt3), on='Date', how='left')
df['GT3_Hits'] = df['GT3_Hits'].replace(0, np.nan).ffill()

```

```
In [ ]: df=df.fillna(value=0)
```

```
In [ ]: earliest_time= df['Date'].min()
```

```

In [ ]: df["day_of_week"] = df['Date'].dt.dayofweek.astype(str) # categories have be strings
df["week_of_year"] = df['Date'].dt.isocalendar().week.astype(str) # categories have be strings
df["month"] = df['Date'].dt.month.astype(str)
df['hour'] = df['Date'].dt.hour.astype(str)
df['day'] = df['Date'].dt.day.astype(str)
df['year'] = df['Date'].dt.year-2006
df['year']=df['year'].astype('int')
#df['minute'] = df['Date'].dt.minute.astype(str).astype("int")
df['time_idx']=df.index
df['hours_from_start'] = (df['Date'] - earliest_time).dt.seconds / 60 / 60 + (df['Date'] - earliest_time).dt.days
df['hours_from_start'] = df['hours_from_start'].astype('int')
df['days_from_start'] = (df['Date'] - earliest_time).dt.days
df['group']=0
df['Global_active_power']=df['Global_active_power'].astype("float64")
df['Voltage']=df['Voltage'].astype("float64")
df['Global_intensity']=df['Global_intensity'].astype("float64")
df['Sub_metering_1']=df['Sub_metering_1'].astype("float64")
df['Sub_metering_2']=df['Sub_metering_2'].astype("float64")
df['Wind_Speed_Norm'] = df['Wind_Speed_Norm'].astype("float64")
df['Prec_Depth_Norm'] = df['Prec_Depth_Norm'].astype("float64")
df['Air_Temp_Norm'] = df['Air_Temp_Norm'].astype("float64")

time_df = df[[
    'Global_active_power'
    , 'Sub_metering_1'
    , 'Sub_metering_2'
    , 'Sub_metering_3'
    , 'CPI'
    , 'CCI'
    , 'GT1_Hits'
    , 'GT2_Hits'
    , 'GT3_Hits'
    , 'Wind_Speed_Norm'
    , 'Prec_Depth_Norm'
    , 'Air_Temp_Norm'
    , 'day_of_week'
    , 'week_of_year'
    , 'year'
    , 'month'
    , 'hour'
    , 'day'
    , 'time_idx'
    , 'days_from_start'
]]

```

```
, 'group'
, 'Date'
]]
df.dtypes
```

```
Out[ ]: Date          datetime64[ns]
Time                object
Global_active_power float64
Global_reactive_power object
Voltage            float64
Global_intensity   float64
Sub_metering_1     float64
Sub_metering_2     float64
Sub_metering_3     float64
CPI                float64
CCI                float64
Wind_Speed_Norm    float64
Prec_Depth_Norm    float64
Air_Temp_Norm      float64
GT1_Hits           float64
GT2_Hits           float64
GT3_Hits           float64
day_of_week        object
week_of_year       object
month             object
hour              object
day               object
year              int64
time_idx           int64
hours_from_start   int64
days_from_start    int64
group              int64
dtype: object
```

```
In [ ]: time_df.isna().values.sum()
```

```
Out[ ]: 0
```

```
In [ ]: time_df.head()
```

```
Out[ ]:   Global_active_power  Sub_metering_1  Sub_metering_2  Sub_metering_3  CPI  CCI  GT1_Hits  GT2_Hits  GT3_Hits
```

0	2.580	0.0	0.0	0.0	72.71	100.0	45.0	41.0	0.0
1	2.460	0.0	0.0	0.0	72.71	100.0	45.0	41.0	0.0
2	2.478	0.0	0.0	0.0	72.71	100.0	45.0	41.0	0.0
3	1.416	0.0	0.0	17.0	72.71	100.0	45.0	41.0	0.0
4	0.442	0.0	0.0	0.0	72.71	100.0	45.0	41.0	0.0

5 rows × 22 columns

Setting up Time Series DataLoaders

```
In [ ]: #Hyperparameters
#batch size=64
#number heads=4, hidden sizes=160, lr=0.001, gr_clip=0.1

max_prediction_length = 4*7*4*6 #sixmonths
max_encoder_length = 4*365 #one year
training_cutoff = time_df["time_idx"].max() - max_prediction_length
```

```

training = TimeSeriesDataSet(
    time_df[Lambda x: x.time_idx <= training_cutoff],
    time_idx="time_idx",
    target="Global_active_power",
    group_ids=['group'],
    min_encoder_length=4*7, #one week
    max_encoder_length=max_encoder_length,
    min_prediction_length=4, #one day
    max_prediction_length=max_prediction_length,
    #static_categoricals=["consumer_id"],
    time_varying_known_categoricals=["day", "day_of_week", "month", 'hour' ], #year to be added
    time_varying_known_reals=["time_idx", 'year'],
    time_varying_unknown_reals=['Sub_metering_1', 'Sub_metering_2', 'Sub_metering_3', 'CPI', 'CCI', 'GT1_Hits'],
    target_normalizer=GroupNormalizer(
        groups=["group"], transformation="softplus"
    ), # we normalize by group
    add_relative_time_idx=True,
    add_target_scales=True,
    add_encoder_length=True,
    #add_nan=True,
    allow_missing_timesteps=True
)

validation = TimeSeriesDataSet.from_dataset(training, time_df, predict=True, stop_randomization=True)

# create dataloaders for our model
batch_size = 64
# if you have a strong GPU, feel free to increase the number of workers
train_dataloader = training.to_dataloader(train=True, batch_size=batch_size, num_workers=10)
val_dataloader = validation.to_dataloader(train=False, batch_size=batch_size * 10, num_workers=10)

```

/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 10 worker processes in total. Our suggested max number of worker in current system is 8, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

warnings.warn(_create_warning_msg(

[I 2023-12-10 07:14:28,911] Trial 3 finished with value: 0.31155362725257874 and parameters: {'gradient_clip_val': 0.8223697202530604, 'hidden_size': 22, 'dropout': 0.2972505585187706, 'hidden_continuous_size': 9, 'attention_head_size': 1, 'learning_rate': 0.014035308702029799}. Best is trial 2 with value: 0.29049456119537354.

Saving the Model

```

In [ ]: #!unzip /content/drive/MyDrive/BDA696Project/Model/model.zip
        #torch.save(best_tft.state_dict(), path)

```

```

In [ ]: best_model_path='/content/drive/MyDrive/BDA696Project/Model/model_saved/lightning_logs/lightning_logs/vers:
        best_tft = TemporalFusionTransformer.load_from_checkpoint(best_model_path)

```

/usr/local/lib/python3.10/dist-packages/lightning/pytorch/utilities/parsing.py:198: Attribute 'loss' is an instance of `nn.Module` and is already saved during checkpointing. It is recommended to ignore them using `self.save_hyperparameters(ignore=['loss'])`.

/usr/local/lib/python3.10/dist-packages/lightning/pytorch/utilities/parsing.py:198: Attribute 'logging_metrics' is an instance of `nn.Module` and is already saved during checkpointing. It is recommended to ignore them using `self.save_hyperparameters(ignore=['logging_metrics'])`.

```

In [ ]: #Take a look at what the raw_predictions variable contains

raw_predictions = best_tft.predict(val_dataloader, mode="raw", return_x=True)
print(raw_predictions._fields)
#('output', 'x', 'index', 'decoder_lengths', 'y')

print('\n')
#print(raw_predictions.output._fields)
# ('prediction',
#  'encoder_attention',
#  'decoder_attention',
#  'static_variables',
#  'encoder_variables',

```



```
# 'decoder_variables',
# 'decoder_lengths',
# 'encoder_lengths')

print('\n')
#print(raw_predictions.output.prediction.shape)
#torch.Size([5, 24, 7])

# We get predictions of 5 time-series for 24 days.
# For each day we get 7 predictions - these are the 7 quantiles:
#[0.02, 0.1, 0.25, 0.5, 0.75, 0.9, 0.98]
# We are mostly interested in the 4th quantile which represents, let's say, the 'median loss'
# fyi, although docs use the term quantiles, the most accurate term are percentiles

# We get predictions of 5 time-series for 24 days.
# For each day we get 7 predictions - these are the 7 quantiles:
#[0.02, 0.1, 0.25, 0.5, 0.75, 0.9, 0.98]
# We are mostly interested in the 4th quantile which represents, let's say, the 'median loss'
# fyi, although docs use the term quantiles, the most accurate term are percentiles
```

```
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will
ll create 10 worker processes in total. Our suggested max number of worker in current system is 8, which is
smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might
get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if n
ecessary.
  warnings.warn(_create_warning_msg(
('output', 'x', 'index', 'decoder_lengths', 'y'))
```

Loss Comparision: TFT VS Baseline

Using 'prediction' mode for calculating loss function

```
In [ ]: actuals = torch.cat([y[0] for x, y in iter(val_dataloader)])
        predictions = best_tft.predict(val_dataloader, mode="prediction").to('cpu')
        baseline_predictions = Baseline().predict(val_dataloader).to('cpu')
```

```
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
```

MAE

```
In [ ]: mae_loss = MAE()  
print(f"TFT: {mae_loss.loss(actuals, predictions).mean(axis = 1).median().item()}")  
print(f"Baseline: {mae_loss.loss(actuals, baseline_predictions).mean(axis = 1).median().item()}")
```

TFT: 0.549065113067627
Baseline: 0.5105714201927185

SMAPE

```
In [ ]: sm = SMAPE()  
print(f"TFT: {sm.loss(actuals, predictions).mean(axis = 1).median().item()}")  
print(f"Baseline: {sm.loss(actuals, baseline_predictions).mean(axis = 1).median().item()}")
```

TFT: 0.6751478910446167
Baseline: 0.729117214679718

MAPE

```
In [ ]: mape = MAPE()  
print(f"TFT: {mape.loss(actuals, predictions).mean(axis = 1).median().item()}")  
print(f"Baseline: {mape.loss(actuals, baseline_predictions).mean(axis = 1).median().item()}")
```

TFT: 0.9968411326408386
Baseline: 1.7605912685394287

Poisson Loss

```
In [ ]: pl = PoissonLoss()  
print(f"TFT: {pl.loss(actuals, predictions).mean(axis = 1).median().item()}")  
print(f"Baseline: {pl.loss(actuals, baseline_predictions).mean(axis = 1).median().item()}")
```

TFT: 3.1876018047332764
Baseline: 3.5891759395599365

Quantile Loss

```
In [ ]: ql = QuantileLoss()  
print(f"TFT: {ql.loss(actuals.view(1, -1), predictions.view(1, -1)).mean().item()}")  
print(f"Baseline: {ql.loss(actuals.view(1, -1), baseline_predictions.view(1, -1)).mean().item()}")
```

TFT: 0.9532883763313293
Baseline: 1.0016515254974365



Observed vs Predicted

Using 'raw' mode


```
In [ ]: import matplotlib.pyplot as plt  
import seaborn as sns  
  
plt.style.use('seaborn-deep')  
fig, ax = plt.subplots(figsize=(10, 4))  
best_tft.plot_prediction(raw_predictions.x, raw_predictions.output, idx=0, add_loss_to_title=True, ax=ax)
```


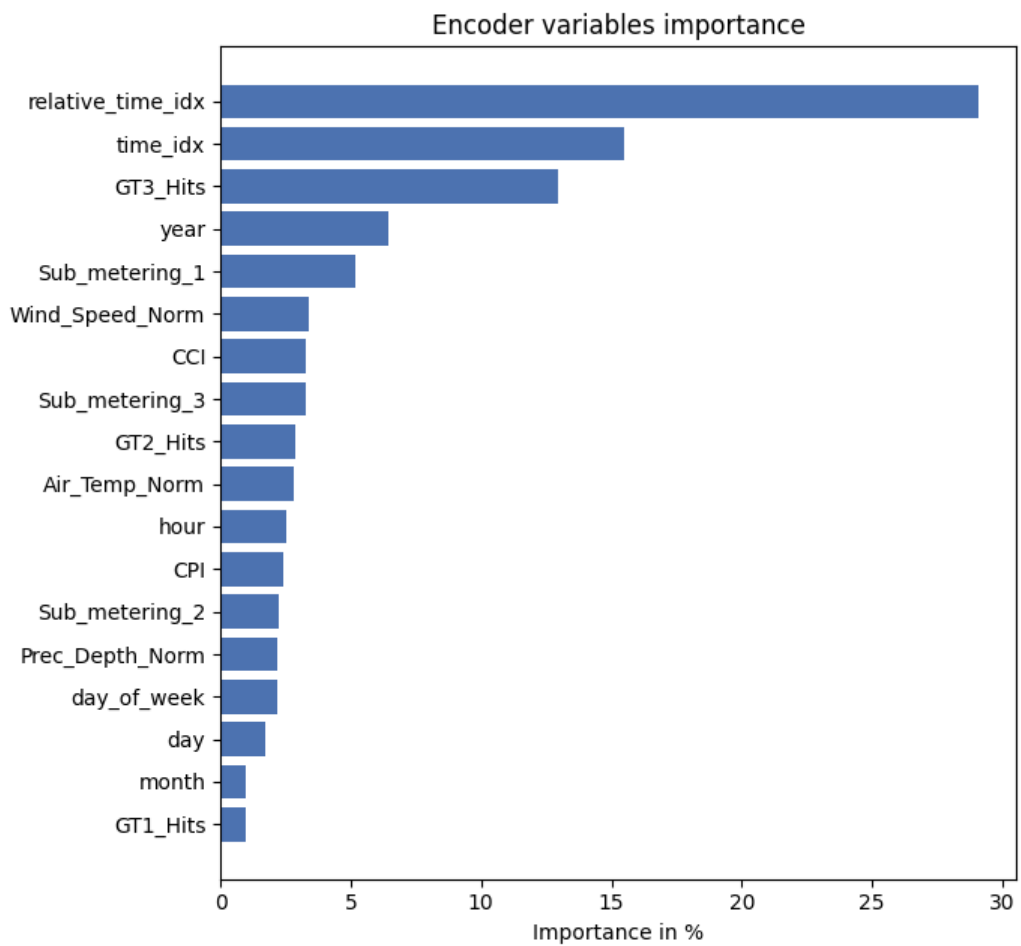

<ipython-input-54-86e257ce8775>:4: MatplotlibDeprecationWarning: The seaborn styles shipped by Matplotlib are deprecated since 3.6, as they no longer correspond to the styles shipped by seaborn. However, they will remain available as 'seaborn-v0_8-'. Alternatively, directly use the seaborn API instead.
plt.style.use('seaborn-deep')

Out[]:



```
In [ ]: interpretation = best_tft.interpret_output(raw_predictions.output, reduction="sum")  
#interpretation  
plt.show(best_tft.plot_interpretation(interpretation))
```



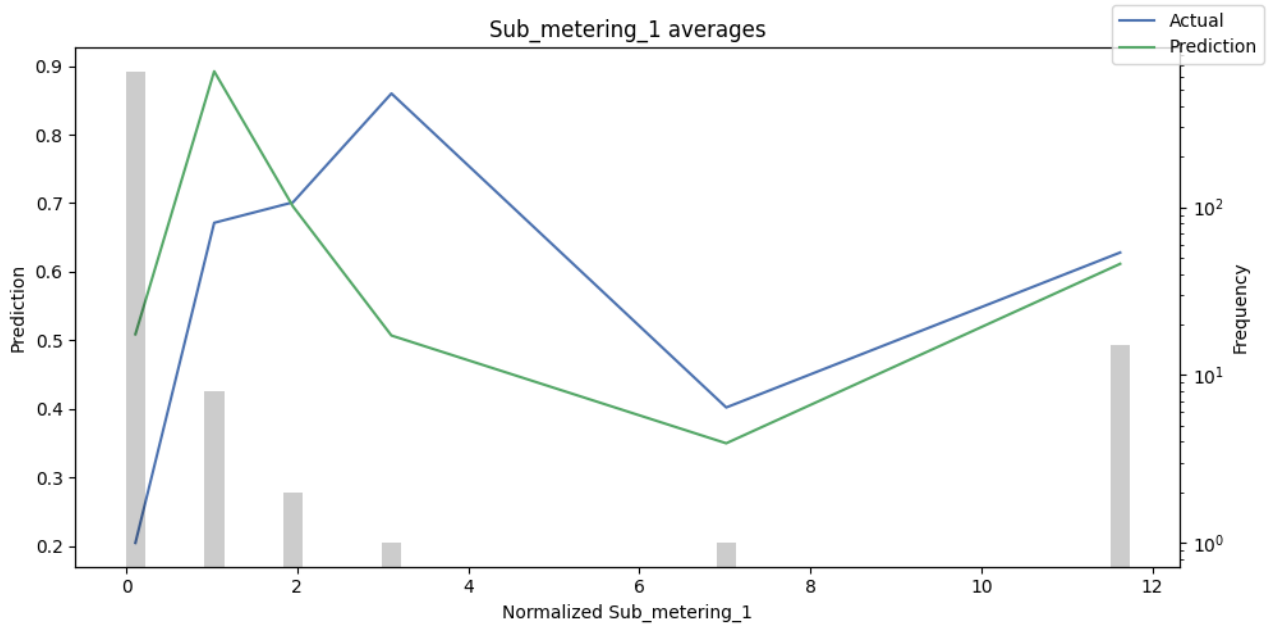



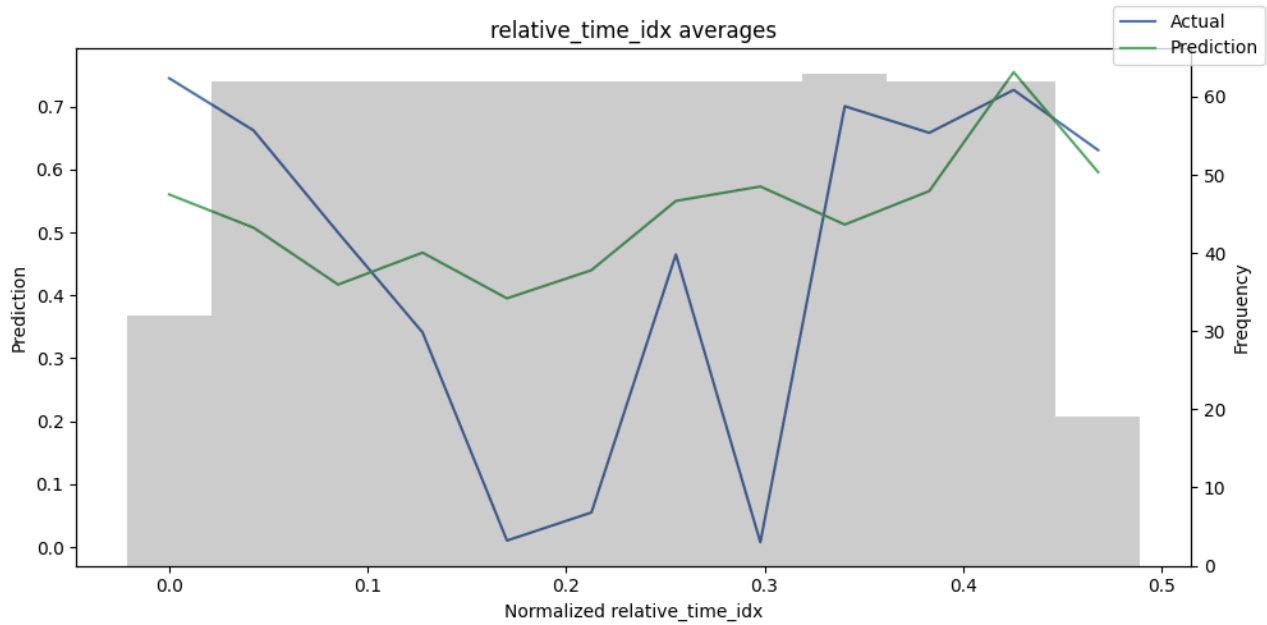
Using 'prediction' mode

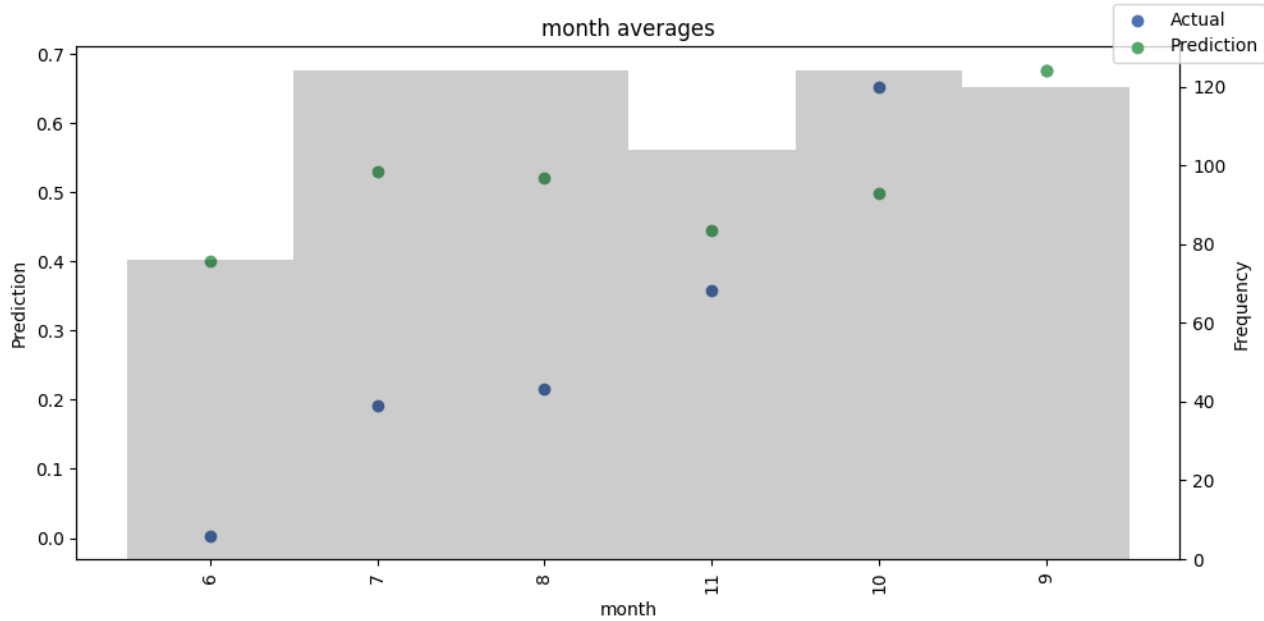
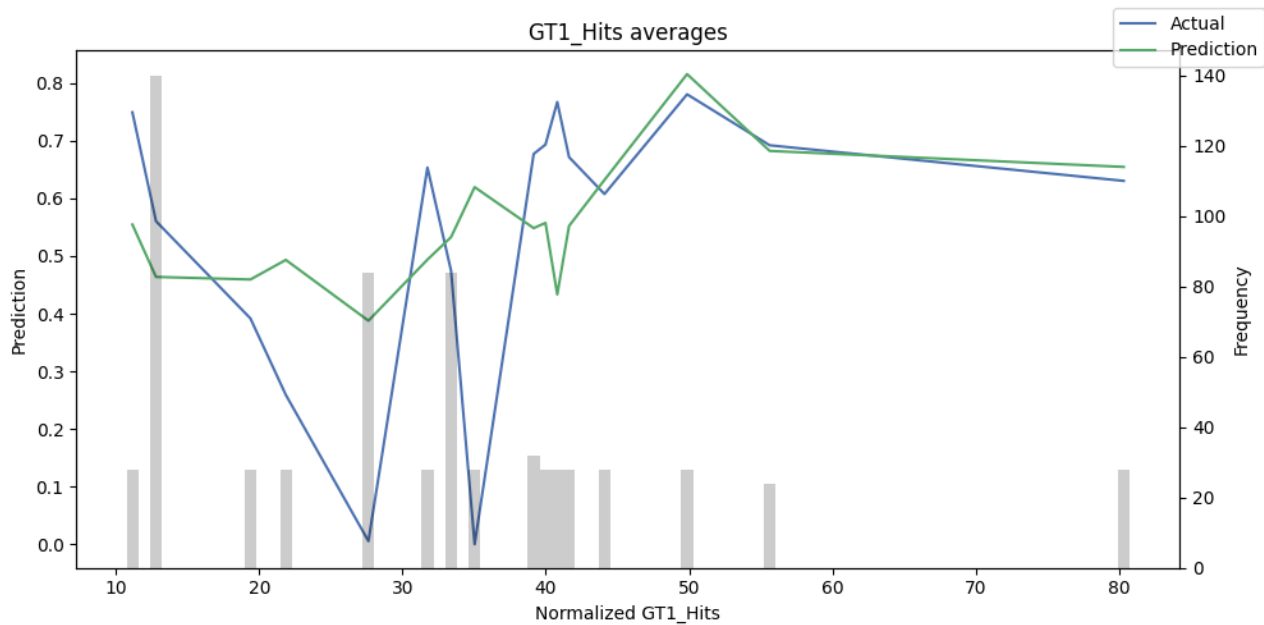
```
In [ ]: predictions=best_tft.predict(val_dataloader,mode='prediction',return_x=True)
```

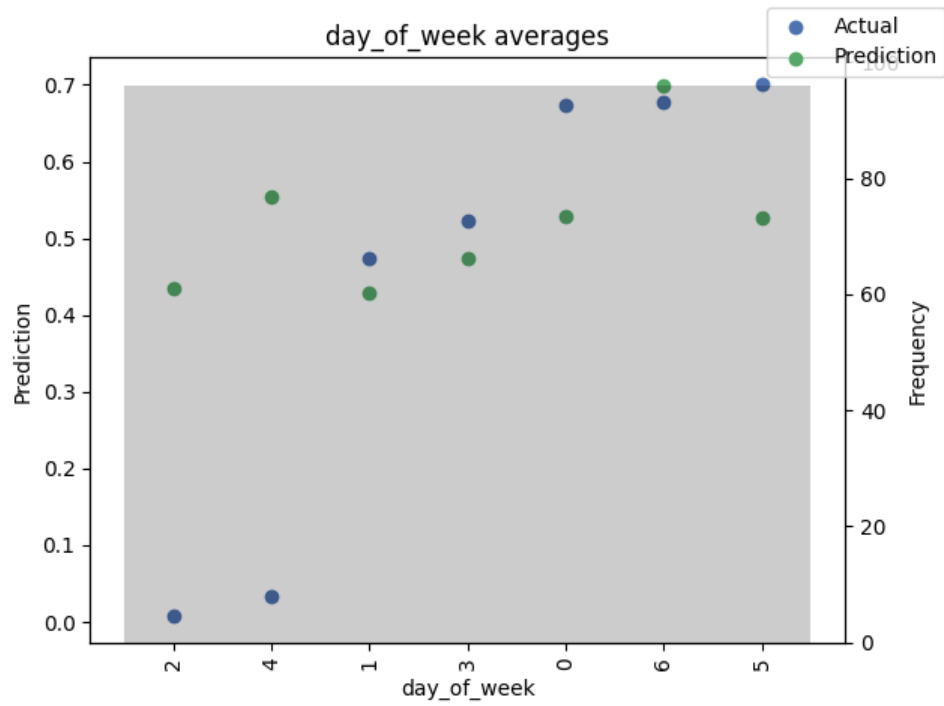
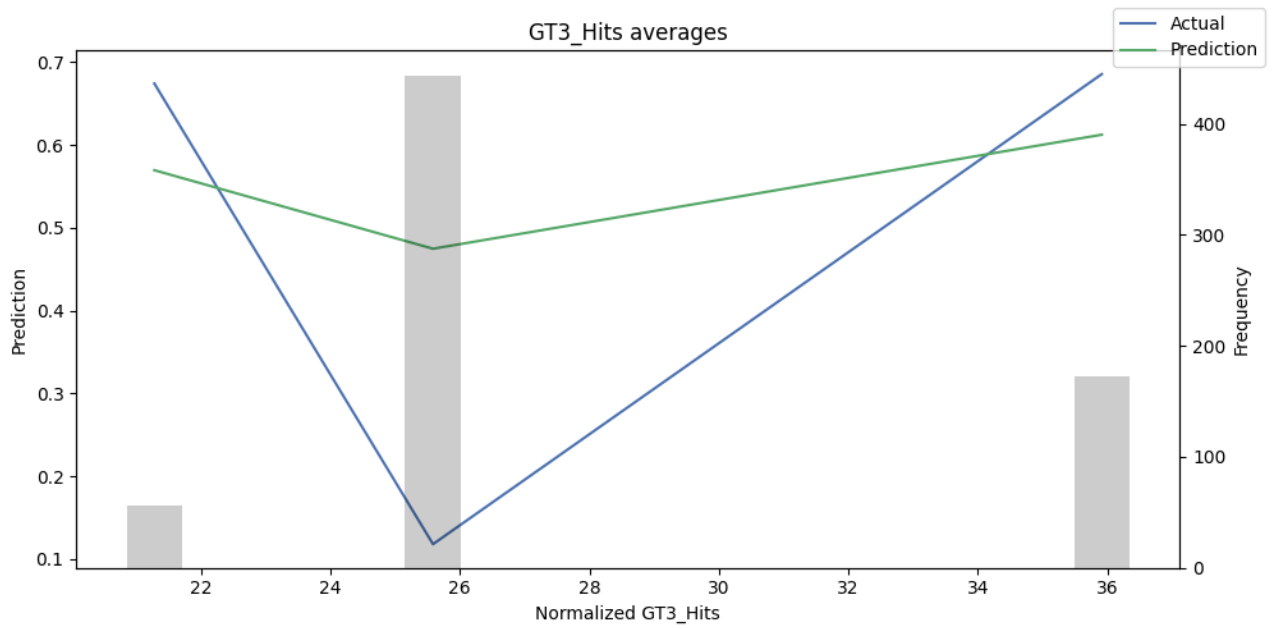
```
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will
ll create 10 worker processes in total. Our suggested max number of worker in current system is 8, which is
smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might
get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if n
ecessary.
  warnings.warn(_create_warning_msg(
```

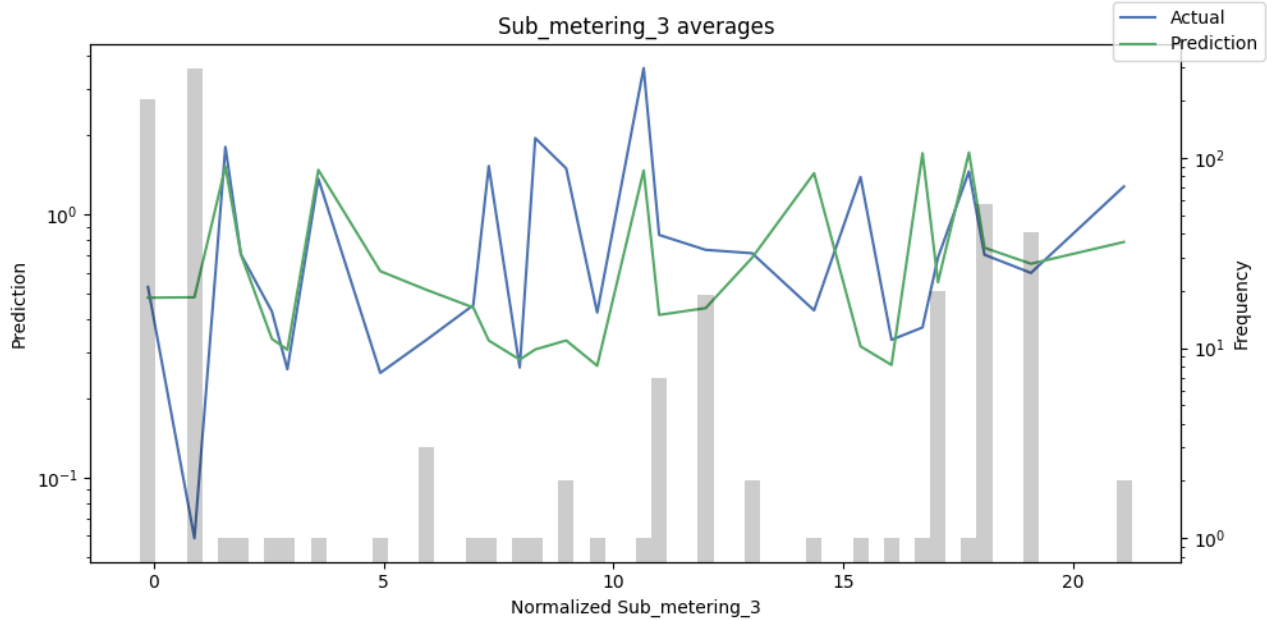
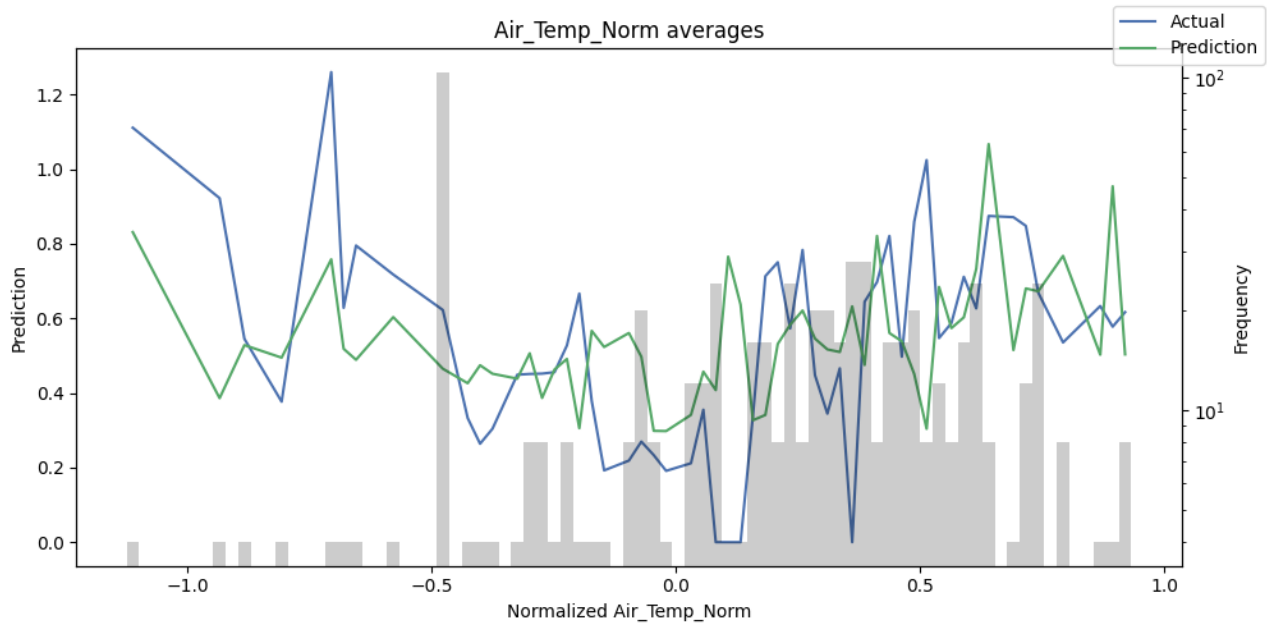
```
In [ ]: predictions_vs_actuals = best_tft.calculate_prediction_actual_by_variable(predictions.x, predictions.output
all_features = list(set(predictions_vs_actuals['support'].keys())-set(['year', 'encoder_length', 'Global_a
for feature in all_features:
    best_tft.plot_prediction_actual_by_variable(predictions_vs_actuals, name=feature);
```


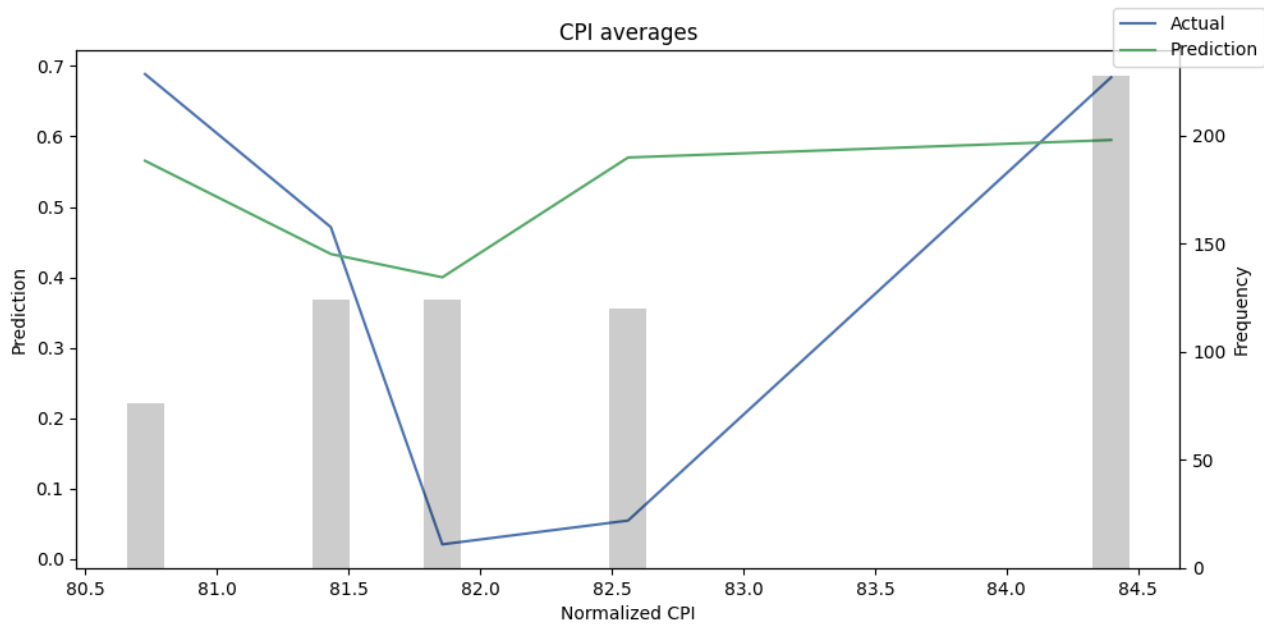



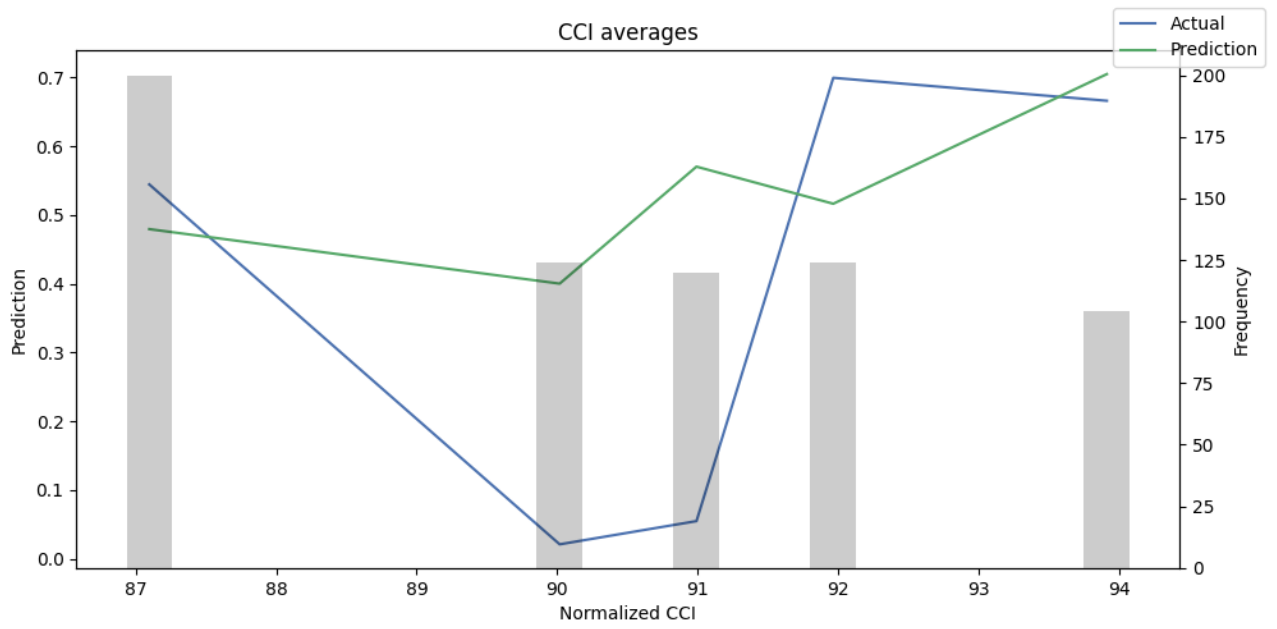






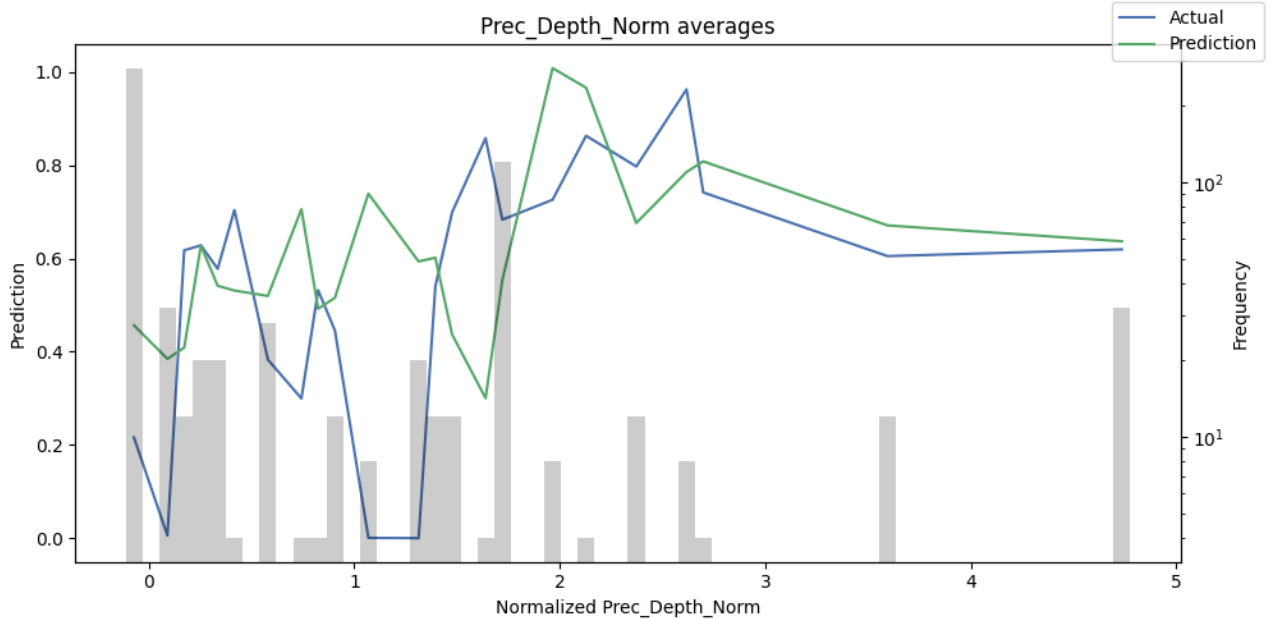
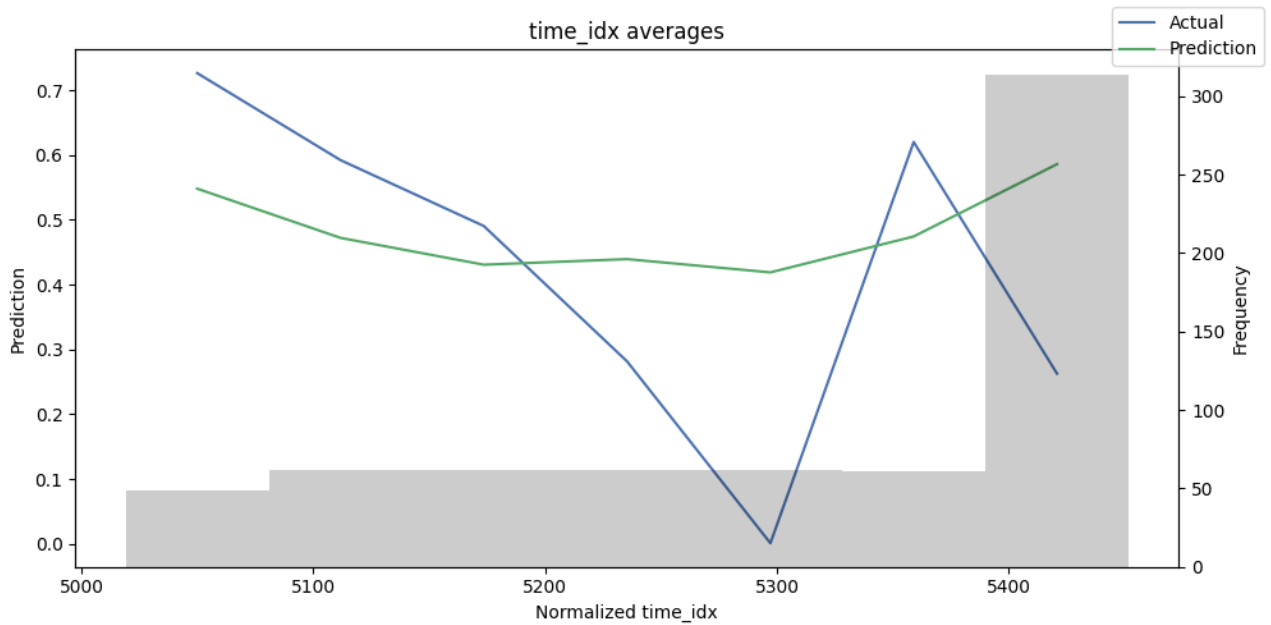












```
In [ ]: dependency = best_tft.predict_dependency(val_data_loader.dataset, 'relative_time_idx', np.linspace(0, 300, 300))
# plotting median and 25% and 75% percentile
agg_dependency = dependency.groupby('relative_time_idx').normalized_prediction.agg(median="median", q25="q25", q75="q75")
ax = agg_dependency.plot(y="median")
ax.fill_between(agg_dependency.index, agg_dependency.q25, agg_dependency.q75, alpha=0.01)

Predict: 0% | 0/30 [00:00<?, ? batches/s]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU cores
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPU cores
INFO: HPU available: False, using: 0 HPU cores
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPU cores
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU cores
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPU cores
INFO: HPU available: False, using: 0 HPU cores
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPU cores
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU cores
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPU cores
INFO: HPU available: False, using: 0 HPU cores
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPU cores
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU cores
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPU cores
INFO: HPU available: False, using: 0 HPU cores
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPU cores
```

```
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero: TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero: IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero: HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero: GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
```

```
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
```



```

es, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument` to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

```

```
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument' to 'num_workers=7' in the 'DataLoader' to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument' to 'num_workers=7' in the 'DataLoader' to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument' to 'num_workers=7' in the 'DataLoader' to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument' to 'num_workers=7' in the 'DataLoader' to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument' to 'num_workers=7' in the 'DataLoader' to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO:lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
```

```
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:GPU available: False, using: 0 GPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPUs
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_data_loader' does not have many workers which may be a bottleneck. Consider increasing the value of the `num_workers` argument to `num_workers=7` in the `DataLoader` to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
```

```
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument to 'num_workers=7' in the 'DataLoader' to improve performance.
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
INFO: GPU available: True (cuda), used: True
INFO: lightning.pytorch.utilities.rank_zero:GPU available: True (cuda), used: True
INFO: TPU available: False, using: 0 TPU cores
INFO: lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO: IPU available: False, using: 0 IPU
INFO: lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO: HPU available: False, using: 0 HPU
INFO: lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPU
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
/usr/local/lib/python3.10/dist-packages/lightning/pytorch/trainer/connectors/data_connector.py:441: The 'predict_dataloader' does not have many workers which may be a bottleneck. Consider increasing the value of the 'num_workers' argument to 'num_workers=7' in the 'DataLoader' to improve performance.
```

Out[]: <matplotlib.collections.PolyCollection at 0x7d0be1430fa0>

